

Algorithms and Data Structures (II)

Course 2,
Gabriel Istrate

March 9, 2023

First of all ...



Example: Selection sort vs. HEAPSORT

Selection SORT

- Elements to sort in a vector.
- Find maximum element.
- Swap it with the last element.
- Proceed recursively.

Complexity

- Finding maximum in a vector: $\Theta(n)$.
- Complexity analysis: $T(n) = T(n - 1) + \theta(n)$.
- Conclusion: $\Theta(n^2)$

Example: Selection sort vs. HEAPSORT

Selection SORT

- Bottleneck in Selection Sort: Find maximum element
- If we could improve finding max

HEAPSORT

- Algorithm: Same idea.
- Bottleneck: FindMax $\theta(n)$. Replace it with $O(1)$ operation (via heaps).
- Complexity $W(n) = \theta(\log n)$ (need also to update heap via HEAPIFY)
- New algorithm: complexity $\theta(n \log(n))$.

Why data structures ?

- Data structures: algorithm development via **primitive operations**.

Modularity in algorithm design

You don't build a house from scratch (bricks, frames, drywalls). **Same with algorithms/code.**

- Easier to solve problem/test solution only once.
- Correctness: easier to check. **easier to update.**

Why data structures ? Performance.

- You google something. Don't want to wait 100 seconds ! Search: fast.
- You play a game. Game engine must quickly retrieve/update objects you see in front of you when you move your viewport.
- Operations: often abstracted from requirements.

Most frequent operations should be fast.

How do you measure performance ?

$O(n \log n)$, $\Theta(\log n)$...

Example (operations from requirements)

$\underbrace{1 \ 2 \ 3} \quad \underbrace{7 \ 5 \ 6 \ 4} \quad \underbrace{8 \ 9 \ 10} \quad \underbrace{14 \ 13 \ 12 \ 11} \dots$

- TCP: basis for much of Internet traffic.
- **Data requirement:** We need to **buffer** a packet that is out-of order.
- We need to **pop** elements that become in-order.
- We need to test emptiness of buffer.
- We need to produce first missing element (ACK).
- Operation performance $O(1)$?

- A data structure is a way to organize and store information
 - ▶ to facilitate access, or for other purposes
- A data structure has an interface consisting of procedures for adding, deleting, accessing, reorganizing, etc.
- A data structure stores data and possibly meta-data
 - ▶ e.g., a heap needs an array A to store the keys, plus a variable $A.\text{heap-size}$ to remember how many elements are in the heap

What are data structures more concretely ?

- data ...
- E.g. complex numbers: two floats.
- ... together with operations one can perform on the data ...
Example: integer + (addition), - (subtraction), \cdot (multiplication).
- ... and performance guarantees.

Note !

How to precisely implement operations is **not** a part of data structure specification. Concepts, not code.

- All DS that share a common structure and expose the same set of operations.
- Predefined data types: array, structures, files.
- Scalar data type: ordering relation exists among elements.
- More complicated: dynamic DS. Lists, circular lists, trees, hash tables, graphs.

C++: Standard template library (STL):

library of container classes, algorithms, and iterators; provides many of the basic algorithms and data structures of computer science

Example: Array data type/Vector

- Ensures **random access** to its elements.
- Complexity $O(1)$.
- Composed of objects of the same type.

Implementations

- `int myarray[10];` One dimensional arrays.
- Multidimensional arrays.
`type name[lim1]...[limn];`
- implementation in C++/STL: `vector`.

Example using vector class

```
#include<vector>
    using namespace std;

    int main(){

        static const int SIZE = 10000;
        vector<int> arr( SIZE );
        arr.append(125);
        ....
    }
```

Vectors the C++/data structures way

- Vector: black-box.
- Random access: `arr[i]` should take $\Theta(1)$ time.
- Black box (class implementation) may implement some other operations, e.g. `append`.

Main point

You didn't implement vector yourself. All you care is **what operations can you execute, and how complex they are.**

This course

Define, implement various "data structures", and use them to get better algorithms.

Some minimal C/C++ recap

You have an entire course for more.

Variables that hold addresses of other variables.

```
int i=15,j, *p,*q;
```

Dynamic memory allocation: `p= new int;`

Assignment: `*p=20;`

Deallocation: `delete p;`

Dangling reference: upon deallocation should
assign `p = 0;`

Pointers and arrays

```
int a[5],*p;
```

```
for(sum=a[0],i=1;i<5;i++)  
    sum += a[i];
```

or

```
for (sum=*a,i=1;i<5;i++)  
    sum += *(a+i);
```

or

```
for(sum=*a,p=a+1;p<a+5;p++)  
    sum += *p;
```

```
p = new int[n];  
delete [] p;
```


Pointers and reference variables

```
int n = 5, *p = &n, &r = n;
```

r is a **reference variable**. Must be initialized in definition as reference to a particular variable.

reference: different name for/constant pointer to variable.

```
cout << n << ' ' << *p << ' ' << r << endl;
```

5 5 5

```
n = 7 (*p = 7, r = 7)
```

```
cout << n << ' ' << *p << ' ' << r << endl;
```

7 7 7

cout: C++ way to print. BEST WAY TO PASS PARAMETER: **const reference variables**;

C++: classes, objects, member functions, oh my !

- in C++: classes - user-defined data types.
- objects: instantiations of classes.
- objects have **behavior**, **member functions**.

Example

- Assume dog is a C++ class.
- Assume Buddy is an "object" of type dog.
- Dogs behavior: bark, member function with no parameters.
- To make Buddy bark: **Buddy.bark()** call member function bark that belongs to Buddy.
- C++: only so-called **public** member functions can be called from outside the class code.


So let's start ...

Today (time permitting):


- Stacks
- Queues
- Dequeues
- Linked Lists
- Skip Lists.

- A Stack is a sequential organization of items in which the last element inserted is the first element removed. They are often referred to as LIFO, which stands for “last in first out.”
- Examples: letter basket, stack of trays, stack of plates.
- **Only element that may be accessed:** the one that was **most recently inserted**.
- There are only two basic operations on stacks, the **push** (insert), and the **pop** (read and delete).


Stacks: Implementation



(a)



(b)



(c)

- (a). Stack representing set $S = \{2, 6, 9, 15\}$.
- (b). After $PUSH(S, 3)$.
- (c). After $POP(S)$.

Operator Precedence Parsing

- We can use the stack class we just defined to parse and evaluate mathematical expressions like:

$$5 * (((9 + 8) * (4 * 6)) + 7)$$

- First, we transform it to **postfix notation**:

$$5 \ 9 \ 8 \ + \ 4 \ 6 \ * \ * \ 7 \ + \ *$$

- Usual form for arithmetic expressions: **infix**. **term1** op **term2**.
- Postfix notation: **term1 term2** op.
- How to convert infix to postfix: **later !**

Evaluating Postfix expressions

Then, the following C++ routine uses a stack to perform this evaluation:

```
1  char c;  
2  Stack acc(50);  
3  int x;  
4  while (cin.get(c))  
5  {  
6  x = 0;  
7  while (c == ' ') cin.get(c);  
8  if (c == '+') x = acc.pop() + acc.pop();  
9  if (c == '*') x = acc.pop() * acc.pop();  
10 while (c ≥ '0' && c ≤ '9')  
11 x = 10*x + (c-'0'); cin.get(c);  
12 acc.push(x);  
13 }  
14 cout << acc.pop();
```

Explanation of code

- We read one character at a time in `c`.
- In `x` we compute the value of the currently evaluated expression.
- After computing it we push the value on the stack - we will need it later.
- When reading an op we take the last two value off the stack and apply the op on them and assign this to `x`.
- When reading a digit we update value of `x` by making the last read digit the least significant one.

- Algorithms (later).
- Recursion removal.
- Reversing things.
- Procedure call and procedure return is similar to matching symbols:
 - ▶ When a procedure returns, it returns to the most recently active procedure.
 - ▶ When a procedure call is made, save current state on the stack. On return, restore the state by popping the stack.
 - ▶ Formal languages: [pushdown automata](#).

- The ubiquitous “first-in first-out” container (FIFO)

- The ubiquitous “first-in first-out” container (FIFO)
- Interface
 - ▶ `Enqueue(Q, x)` adds element x at the back of queue Q
 - ▶ `Dequeue(Q)` extracts the element at the head of queue Q

- The ubiquitous “first-in first-out” container (FIFO)
- Interface
 - ▶ `Enqueue(Q, x)` adds element x at the back of queue Q
 - ▶ `Dequeue(Q)` extracts the element at the head of queue Q
- Implementation
 - ▶ Q is an array of fixed length $Q.length$
 - ★ i.e., Q holds at most $Q.length$ elements
 - ★ enqueueing more than Q elements causes an “overflow” error
 - ▶ $Q.head$ is the position of the “head” of the queue
 - ▶ $Q.tail$ is the first empty position at the tail of the queue


```
Enqueue(Q,x)
1  if Q.queue-full
2      error "overflow"
3  else Q[Q.tail] = x
4      if Q.tail < Q.length
5          Q.tail = Q.tail + 1
6      else Q.tail = 1
7      if Q.tail == Q.head
8          Q.queue-full = true
9      Q.queue-empty = false
```



Enqueue(Q,x)

```


1  if Q.queue-full
2      error "overflow"
3  else Q[Q.tail] = x
4      if Q.tail < Q.length
5          Q.tail = Q.tail + 1
6      else Q.tail = 1
7      if Q.tail == Q.head
8          Q.queue-full = true
9      Q.queue-empty = false
    
```



Enqueue(Q,x)

```


1  if Q.queue-full
2      error "overflow"
3  else Q[Q.tail] = x
4      if Q.tail < Q.length
5          Q.tail = Q.tail + 1
6      else Q.tail = 1
7      if Q.tail == Q.head
8          Q.queue-full = true
9      Q.queue-empty = false
    
```



Enqueue(Q,x)


```

1  if Q.queue-full
2      error "overflow"
3  else Q[Q.tail] = x
4      if Q.tail < Q.length
5          Q.tail = Q.tail + 1
6      else Q.tail = 1
7      if Q.tail == Q.head
8          Q.queue-full = true
9      Q.queue-empty = false
    
```




```


Enqueue(Q,x)
1  if Q.queue-full
2      error "overflow"
3  else Q[Q.tail] = x
4      if Q.tail < Q.length
5          Q.tail = Q.tail + 1
6      else Q.tail = 1
7      if Q.tail == Q.head
8          Q.queue-full = true
9      Q.queue-empty = false
    
```



Enqueue(Q,x)

```


1  if Q.queue-full
2      error "overflow"
3  else Q[Q.tail] = x
4      if Q.tail < Q.length
5          Q.tail = Q.tail + 1
6      else Q.tail = 1
7      if Q.tail == Q.head
8          Q.queue-full = true
9      Q.queue-empty = false
    
```



Enqueue(Q,x)

```


1  if Q.queue-full
2      error "overflow"
3  else Q[Q.tail] = x
4      if Q.tail < Q.length
5          Q.tail = Q.tail + 1
6      else Q.tail = 1
7      if Q.tail == Q.head
8          Q.queue-full = true
9      Q.queue-empty = false
    
```



Enqueue(Q,x)

```

1  if Q.queue-full
2      error "overflow"
3  else Q[Q.tail] = x
4      if Q.tail < Q.length
5          Q.tail = Q.tail + 1
6      else Q.tail = 1
7      if Q.tail == Q.head
8          Q.queue-full = true
9      Q.queue-empty = false
    
```



Deque(Q)

```


1  if Q.queue-empty
2      error "underflow"
3  else x = Q[Q.head]
4      if Q.head < Q.length
5          Q.head = Q.head + 1
6      else Q.head = 1
7      if Q.tail == Q.head
8          Q.queue-empty = true
9      Q.queue-full = false
10     return x
    
```



Dequeue(Q)

```


1  if Q.queue-empty
2      error "underflow"
3  else x = Q[Q.head]
4      if Q.head < Q.length
5          Q.head = Q.head + 1
6      else Q.head = 1
7      if Q.tail == Q.head
8          Q.queue-empty = true
9      Q.queue-full = false
10     return x
    
```



Deque(Q)

```


1  if Q.queue-empty
2      error "underflow"
3  else x = Q[Q.head]
4      if Q.head < Q.length
5          Q.head = Q.head + 1
6      else Q.head = 1
7      if Q.tail == Q.head
8          Q.queue-empty = true
9      Q.queue-full = false
10     return x
    
```



Deque(Q)

```

1  if Q.queue-empty
2      error "underflow"
3  else x = Q[Q.head]
4      if Q.head < Q.length
5          Q.head = Q.head + 1
6      else Q.head = 1
7      if Q.tail == Q.head
8          Q.queue-empty = true
9      Q.queue-full = false
10     return x
    
```



Dequeue(Q)

```

1  if Q.queue-empty
2      error "underflow"
3  else x = Q[Q.head]
4      if Q.head < Q.length
5          Q.head = Q.head + 1
6      else Q.head = 1
7      if Q.tail == Q.head
8          Q.queue-empty = true
9      Q.queue-full = false
10     return x
    
```

Q.head



Q.tail

Applications of Queues

- Scheduling (disk, CPU)
- Used by operating systems to handle congestion.
- Algorithms (we'll see): breadth-first search.

Stacks, Queues: Scorecard

Algorithm	Complexity
Stack-Empty	$O(1)$ ✓
Push	$O(1)$ ✓
Pop	$O(1)$ ✓
Enqueue	$O(1)$ ✓
Dequeue	$O(1)$ ✓
Restrictions:	LIFO/FIFO orders only. ✗

- Like queues but can enqueue/dequeue at both ends.
- Can modify the code for queues, add two more procedure.
- **do it !**
- Complexity scorecard: similar to queues.

Major problem this semester:

Represent a set S whose elements may vary through time. May want to perform some of:

- $\text{INSERT}(S, x)$
- $\text{DELETE}(S, x)$
- $\text{SEARCH}(S, x)$. Result YES/NO. Better: handle for x , if found.
- $\text{MIN}(S)$
- $\text{MAX}(S)$
- $\text{SUCC}(S, x)$, $\text{PRED}(S, x)$

Example: stacks/queues

- Stacks: dynamic sets with LIFO order.
- Queues: dynamic sets with FIFO order.

- A dictionary is an abstract data structure that represents a set of elements (or keys)
 - ▶ a dynamic set

- A dictionary is an abstract data structure that represents a set of elements (or keys)
 - ▶ a dynamic set
- Interface (generic interface)
 - ▶ $\text{Insert}(D, k)$ adds a key k to the dictionary D
 - ▶ $\text{Delete}(D, k)$ removes key k from D
 - ▶ $\text{Search}(D, k)$ tells whether D contains a key k

- A dictionary is an abstract data structure that represents a set of elements (or keys)
 - ▶ a dynamic set
- Interface (generic interface)
 - ▶ $\text{Insert}(D, k)$ adds a key k to the dictionary D
 - ▶ $\text{Delete}(D, k)$ removes key k from D
 - ▶ $\text{Search}(D, k)$ tells whether D contains a key k
- Implementation
 - ▶ many (concrete) data structures

- A dictionary is an abstract data structure that represents a set of elements (or keys)
 - ▶ a dynamic set
- Interface (generic interface)
 - ▶ `Insert(D, k)` adds a key k to the dictionary D
 - ▶ `Delete(D, k)` removes key k from D
 - ▶ `Search(D, k)` tells whether D contains a key k
- Implementation
 - ▶ many (concrete) data structures
 - ▶ we'll see: hash tables

Direct-Address Table

- A direct-address table implements a dictionary

Direct-Address Table

- A direct-address table implements a dictionary
- The universe of keys is $U = \{1, 2, \dots, M\}$

Direct-Address Table

- A direct-address table implements a dictionary
- The universe of keys is $U = \{1, 2, \dots, M\}$
- Implementation
 - ▶ an array T of size M
 - ▶ each key has its own position in T

Direct-Address Table

- A direct-address table implements a dictionary
- The universe of keys is $U = \{1, 2, \dots, M\}$
- Implementation
 - ▶ an array T of size M
 - ▶ each key has its own position in T

Direct-Address-Insert(T, k)

```
1   $T[k] = \text{true}$ 
```

Direct-Address-Delete(T, k)

```
1   $T[k] = \text{false}$ 
```

Direct-Address-Search(T, k)

```
1  return  $T[k]$ 
```

Direct-Address Table (2)

- Complexity

Direct-Address Table (2)

- Complexity

All direct-address table operations are $O(1)$ ✓

Direct-Address Table (2)

- Complexity

All direct-address table operations are $O(1)$ ✓

So why isn't every set implemented with a direct-address table?

Direct-Address Table (2)

- Complexity

All direct-address table operations are $O(1)$ ✓

So why isn't every set implemented with a direct-address table?

- Space complexity is $\Theta(|U|) \times$

- ▶ $|U|$ is typically a very large number— U is the universe of keys!
- ▶ the represented set is typically much smaller than $|U|$
 - ★ i.e., a direct-address table usually wastes a lot of space

Direct-Address Table (2)

- Complexity

All direct-address table operations are $O(1)$ ✓

So why isn't every set implemented with a direct-address table?

- Space complexity is $\Theta(|U|) \times$

- ▶ $|U|$ is typically a very large number— U is the universe of keys!
- ▶ the represented set is typically much smaller than $|U|$
 - ★ i.e., a direct-address table usually wastes a lot of space

- Want: the benefits of a direct-address table but with a table of reasonable size.

Direct Access Tables: Scorecard

Algorithm	Complexity
INSERT	$O(1)\checkmark$
DELETE	$O(1)\checkmark$
SEARCH	$O(1)\checkmark$
MEMORY:	$\theta(M)\times$

- Interface

- ▶ $\text{List-Insert}(L, x)$ adds element x at beginning of a list L
- ▶ $\text{List-Delete}(L, x)$ removes element x from a list L
- ▶ $\text{List-Search}(L, k)$ finds an element whose key is k in a list L


- Interface

- ▶ `List-Insert(L, x)` adds element `x` at beginning of a list `L`
- ▶ `List-Delete(L, x)` removes element `x` from a list `L`
- ▶ `List-Search(L, k)` finds an element whose key is `k` in a list `L`

- Implementation

- ▶ a doubly-linked list
- ▶ each element `x`: two “links” `x.prev` and `x.next` to the previous and next elements, respectively
- ▶ each element `x`: key `x.key`

Linked List: Implementation



- (a). Linked list representing set $S = \{1, 4, 9, 16\}$.
- (b). After $LIST-INSERT(S, 25)$.
- (c). After $LIST-DELETE(S, 4)$.

Linked List: Implementation

List-Init(L)

```
1  L.head = NIL
```

List-Insert(L, x)

```
1  x.next = L.head
2  if L.head  $\neq$  NIL
3      L.head.prev = x
4      L.head = x
5      x.prev = NIL
```

List-Search(L, k)

```
1  x = L.head.next
2  while x  $\neq$  NIL  $\wedge$  x.key  $\neq$  k
3      x = x.next
4  return x
```


Linked List: Implementation (II)

List-Delete(L, x)

```
1  if x.prev  $\neq$  NIL
2      x.prev.next = x.next
3  else L.head = x.next
4  if x.next  $\neq$  NIL
5      x.next.prev = x.prev
```

Linked List With a “Sentinel”

- instead of NIL sometimes convenient to have a dummy “sentinel” element $L.nil$
- Simplifies LIST-DELETE .
- Adds more memory \times .

Linked List With a “Sentinel”

List-Init(L)

```
1  L.nil.prev = L.nil  
2  L.nil.next = L.nil
```

List-Insert(L, x)

```
1  x.next = L.nil.next  
2  L.nil.next.prev = x  
3  L.nil.next = x  
4  x.prev = L.nil
```

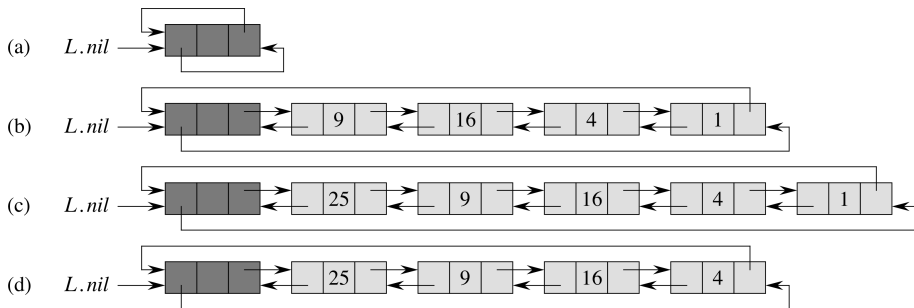
List-Search(L, k)

```
1  x = L.nil.next  
2  while  $x \neq L.nil \wedge x.key \neq k$   
3       $x = x.next$   
4  return x
```

Linked Lists: Observations on Implementation

- Insert: at the head of the list.
- Possible: insert arbitrary position.

Circular Linked Lists



- Can use nil sentinel as head of the list.
- (a): empty circular list.
- (b): Linked list representing set $S = \{1, 4, 9, 16\}$.
- (c): After $LIST-INSERT(S, 25)$.
- (d): After $LIST-DELETE(S, 4)$.

Linked Lists: Scorecard

Linked Lists: Scorecard

Algorithm	Complexity
List-Insert	

Linked Lists: Scorecard

Algorithm	Complexity
List-Insert	$O(1)$ ✓
List-Delete (with pointer)	

Linked Lists: Scorecard

Algorithm	Complexity
List-Insert	$O(1)$ ✓
List-Delete (with pointer)	$O(1)$ ✓
List-Search	

Linked Lists: Scorecard

Algorithm	Complexity
List-Insert	$O(1)$ ✓
List-Delete (with pointer)	$O(1)$ ✓
List-Search	$\Theta(n)$ ✗

Linked Lists: to conclude

- Can reimplement Stacks/Queues using Linked Lists.
- Implementation with pointers: **will not pass the class if you don't know it !**

Advanced topic - Skip lists

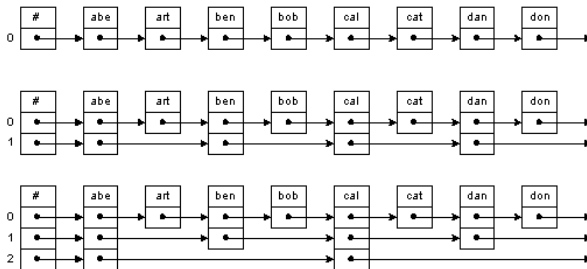
Caution

Topic not in Cormen. See Drozdek for details/C++ implementation.

- Problem with linked list: **search is slow !...** even when elements sorted.
- Solution: **lists of ordered elements** that allow skipping some elements to speed up search.
- **Skip lists**: variant of ordered linked lists that makes such search possible.

More advanced data structure (W. Pugh "Skip lists: a Probabilistic Alternative to Balanced Trees", Communication of the ACM 33(1990), pp. 668-676.) **If anyone curious/interested in data structures/algorithms, can give paper to read; taste how a research article looks like.**

Skip lists



Too theoretical ?

Where does this ever get applied ?

...

Skip lists in real life

According to Wikipedia:

- [MemSQL](#) - skip lists as prime indexing structure for its database technology.
- [Cyrus IMAP server](#) - "skiplist" backend DB implementation
- [Lucene](#) uses skip lists to search delta-encoded posting lists in logarithmic time.
- [QMap](#) (up to Qt 4) template class of Qt that provides a dictionary.
- [Redis](#), ANSI-C open-source persistent key/value store for Posix systems, skip lists in implementation of ordered sets.
- [nessDB](#), a very fast key-value embedded Database Storage Engine.
- [skipdb](#): open-source DB format using ordered key/value pairs.
- [ConcurrentSkipListSet](#) and [ConcurrentSkipListMap](#) in the [Java 1.6 API](#).

Skip lists in real life (II)

According to Wikipedia:

- **Speed Tables**: fast key-value datastore for Tcl that use skiplists for indexes and lockless shared memory.
- **leveldb**, a fast key-value storage library written at Google that provides an ordered mapping from string keys to string values
- **MuQSS** Scheduler for the Linux kernel uses skip lists
- **SkipMap** uses skip lists as base data structure to build a more complex 3D Sparse Grid for Robot Mapping systems.

Skip lists: implementation

What we want

$k = 1, \dots, \lfloor \log_2(n) \rfloor, 1 \leq i \leq \lfloor n/2^{k-1} \rfloor - 1.$

- Item $2^{k-1} \cdot i$ points to item $2^{k-1} \cdot (i + 1).$
 - every second node points to positions two node ahead,
 - every fourth node points to positions four nodes ahead,
 - every eighth node points to positions eighth nodes ahead,
 -, and so on.
-
- Different number of pointers in different nodes in the list !
 - half the nodes only one pointer.
 - a quarter of the nodes two pointers,
 - an eighth of the nodes four pointers,
 -, and so on.
 - $n \log_2(n)/2$ pointers.

Search Algorithm

- ① First follow pointers on the highest level until a larger element is found or the list is exhausted.
- ② If a larger element is found, restart search from its predecessor, this time on a lower level.
- ③ Continue doing this until element found, or you reach the first level and a larger element or the end of the list.

Inserting and deleting nodes

Major problem

- When inserting/deleting a node, pointers of prev/next nodes have to be restructured.
- Solution: rather than equal spacing, **random spacing** on a level.
- Invariant: **Number of nodes on each level: equal, in expectation to what it would be under equal spacing**

Principle

If you're traveling 10 meters in 10 steps, a step is **on average** one meter.

Inserting and deleting nodes (II)

- Level numbering: start with zero.
- New node inserted: probability $1/2$ on first level, $1/4$ second level, $1/8$ third level, ..., etc.
- Function chooseLevel: chooses randomly the level of the new node.
- Generate random number. If in $[0, 1/2]$ level 1, $[1/2, 3/4]$ level 2, etc.
- To delete node: have to update all links.

Computing the i 'th element faster than in $O(i)$

- If we record “step sizes” in our lists we can even mimic indexing !
- Start on highest level.
- If step too big, restart search from predecessor, this time on a lower level.
- Continue doing this until element found.

Update “step sizes” by insertion/deletion

Easy if you have doubly linked lists.

- On deletion: $\text{pred}[i].\text{size}+ = \text{deleted.size}$ on all levels i .
- On insertion: Simply keep track of predecessors and index of the inserted sequence.

Skip Lists: Scorecard

Method	Average	Worst-Case
SPACE:	$O(n)$	$O(n \log(n))$
✓		
SEARCH:	$O(\log(n))$	$O(n)$
✓		
INSERT:	$O(\log(n))$	$O(n)$
✓		
DELETE:	$O(\log(n))$	$O(n)$
✓		

- quite practical ! ✓
- Probabilistic, worst-case still bad. ×
- Not completely easy to implement. ×.

Compared to what ?

Binary search trees. Will learn about them later.